FlashOCC occupancy branch: image encoder → LSS BEV → BEVOCCHead2D

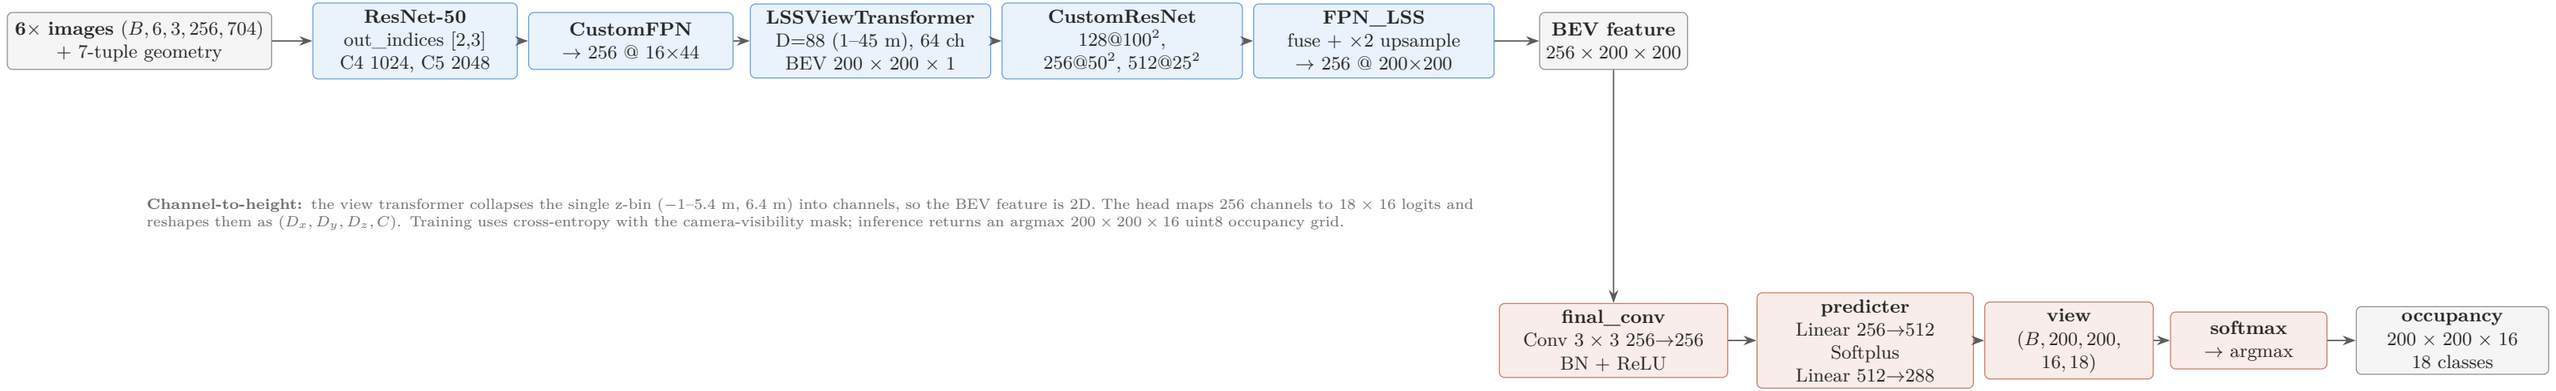

Channel-to-height: the view transformer collapses the single z-bin (−1–5.4 m, 6.4 m) into channels, so the BEV feature is 2D. The head maps 256 channels to 18 × 16 logits and reshapes them as (D_x, D_y, D_z, C). Training uses cross-entropy with the camera-visibility mask; inference returns an argmax 200 × 200 × 16 uint8 occupancy grid.

Occ config: grid $x, y \in [-40, 40]$ @ 0.4 m (200×200), $z \in [-1, 5.4]$ (one bin + 16-way channel folding); depth 1–45 m @ 0.5 m (D=88). The encoder is camera-only (no LiDAR); the 7-tuple input bundles images with sensor2ego / ego2global / intrinsics / post-rots / bda.