TransFusionHead: heatmap-guided proposals + transformer decoder

Heatmap branch: dense per-class heatmap → 3×3 NMS → top-200 proposals. Query features are gathered from the 128-ch BEV map and augmented with a class embedding; the decoder refines them with self/cross attention. Score = sigmoid(heatmap) × query_heatmap_score × one-hot; the bbox coder converts center/height/dim/rot/vel to gravity-center boxes.

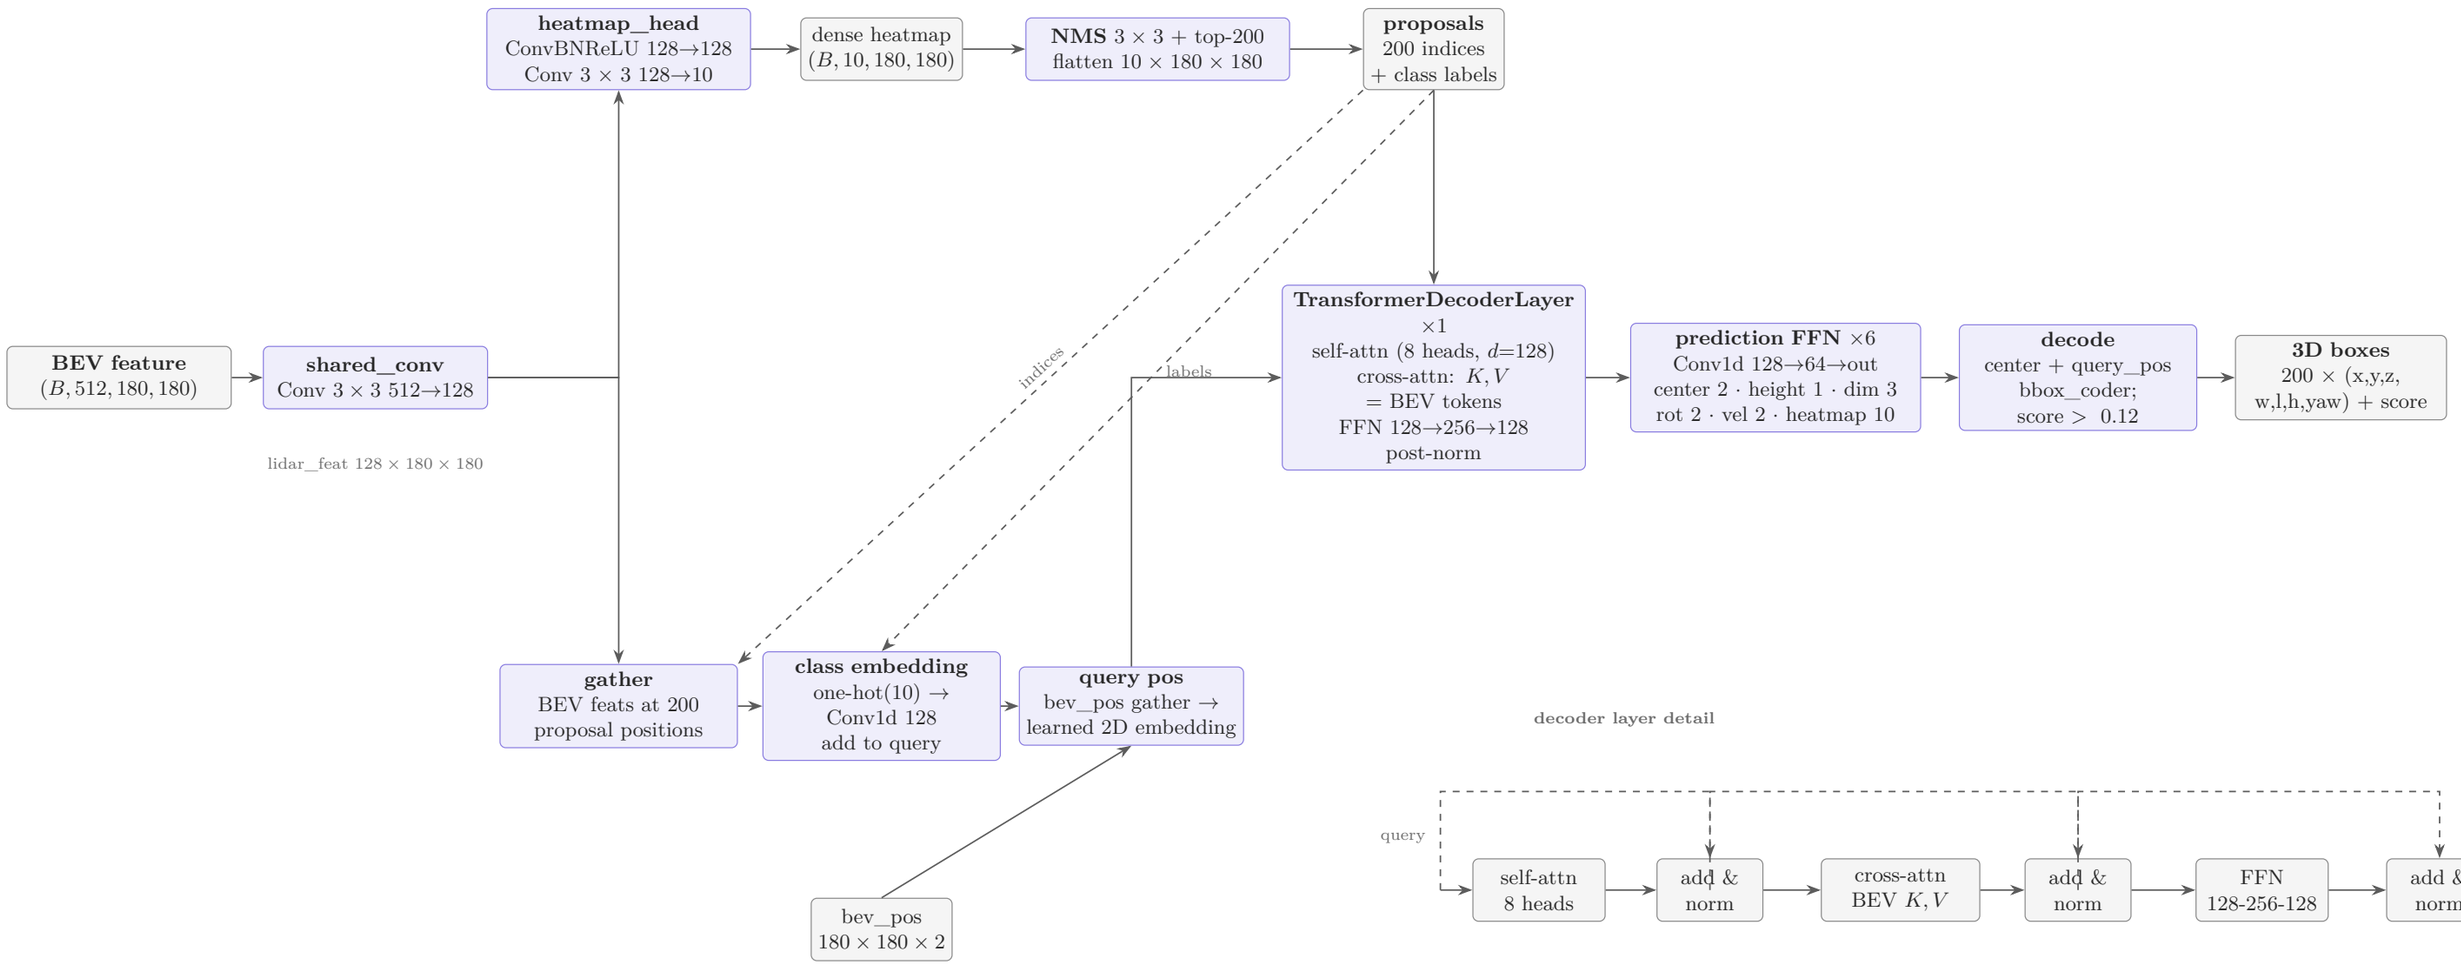

Queries: 200; keys/values: 180×180 BEV tokens with learned 2D position embeddings. Query position is updated from the predicted center for the next layer (1 layer in this config).