View transform: DepthLSSTransform (LSS + LiDAR-conditioned depth)

Depth branch predicts a 118-bin depth distribution per camera pixel and multiplies it with the 80-ch context to build the 5D frustum feature. The geometry branch lifts the frustum to LiDAR coordinates; bev_pool splats both into the BEV grid.

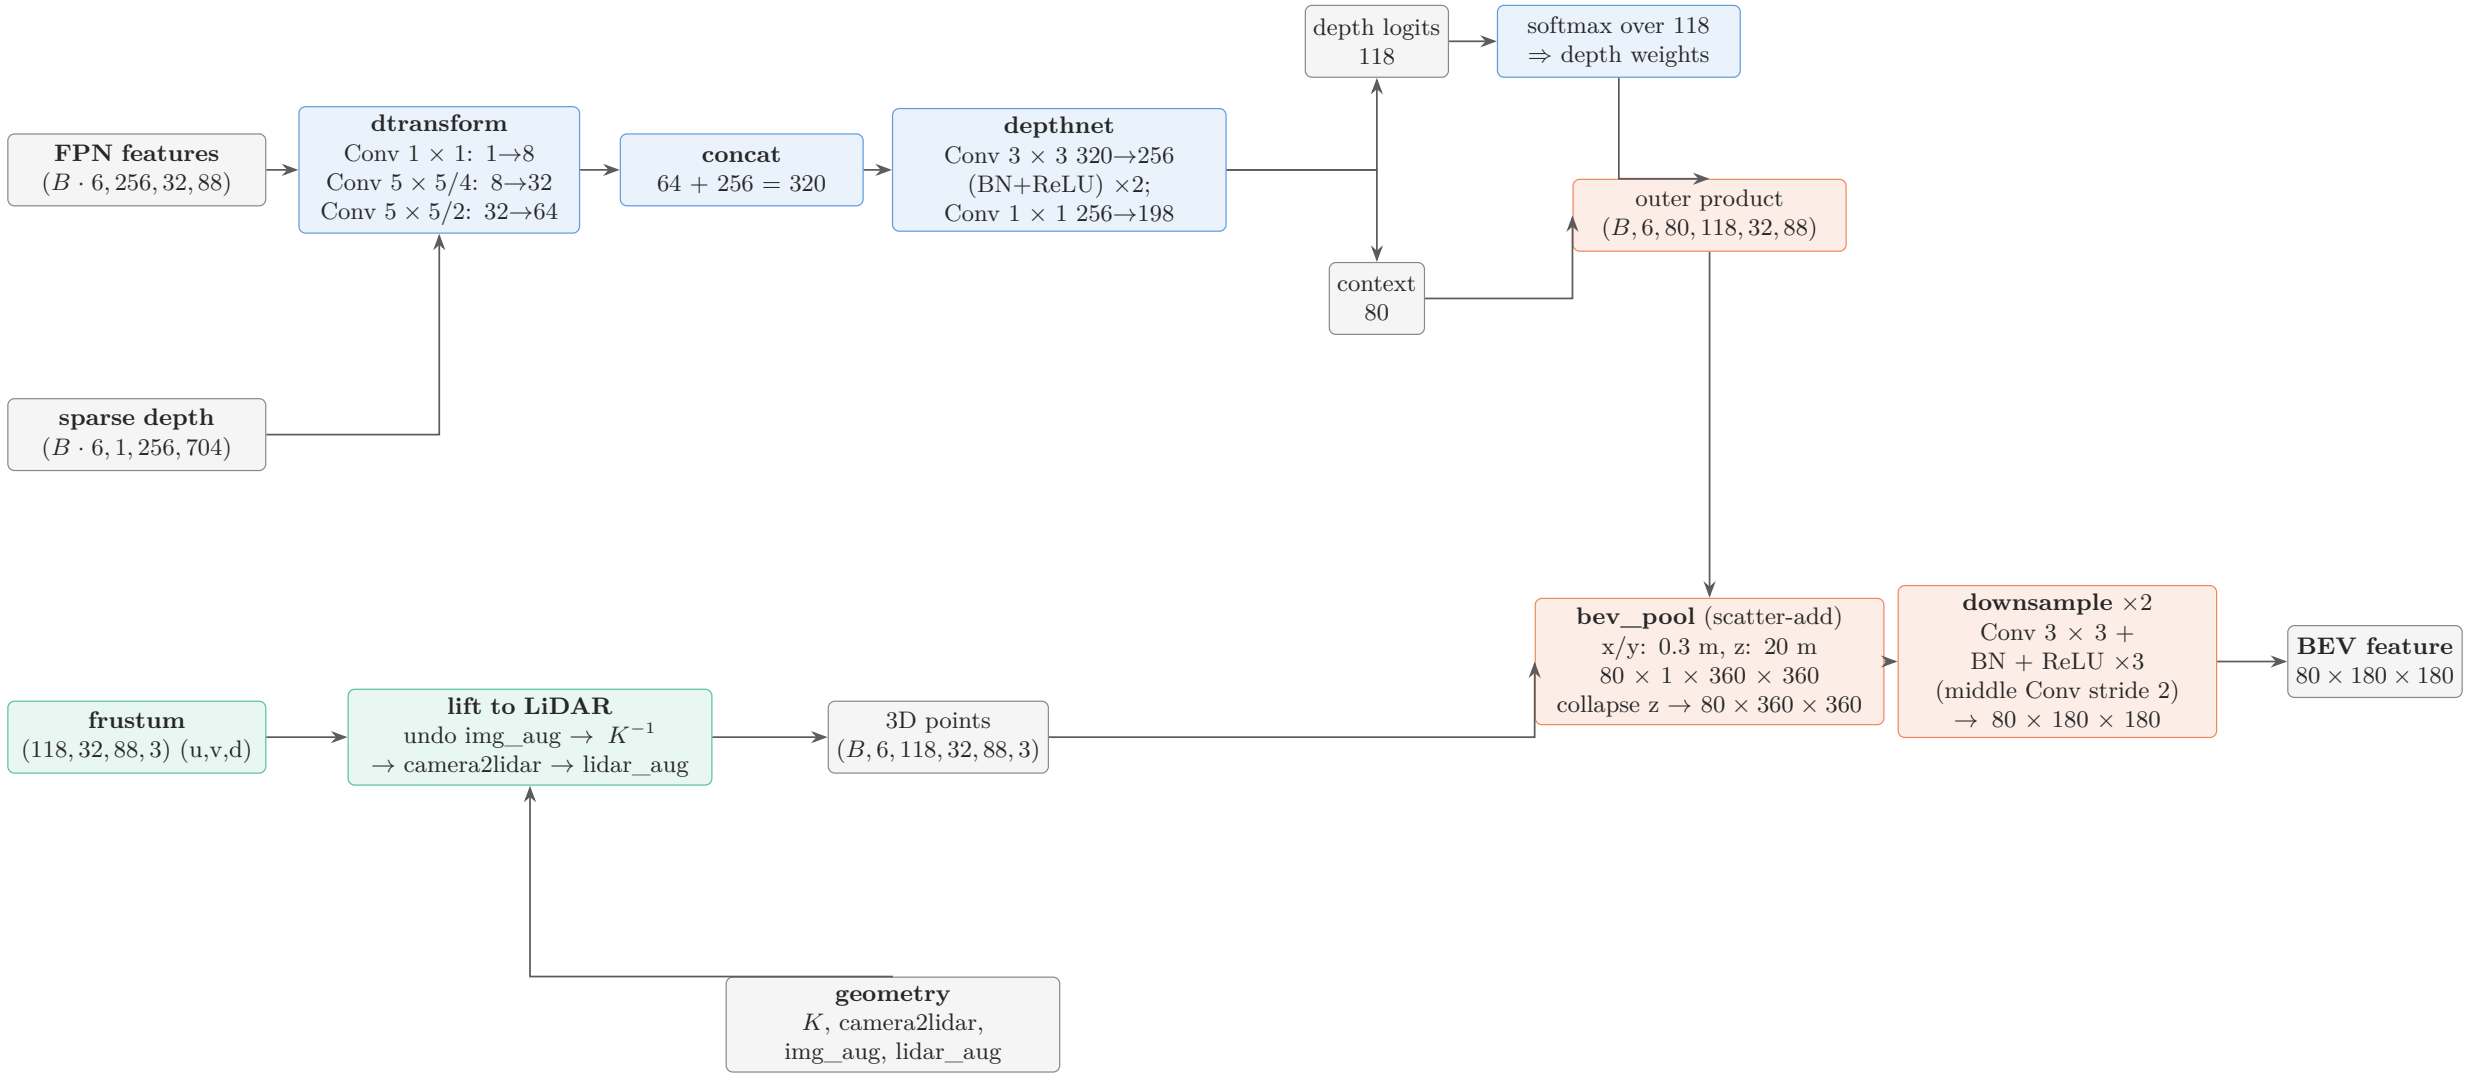

Det config (DepthLSS): $D = 118$ bins, 1–60 m @ 0.5 m; BEV ±54 m, x/y 0.3 m $\to$ 360×360, z 1 bin; downsample 2 $\to$ 180×180. Seg config uses plain LSS (no depth branch), ±51.2 m @ 0.4 m. FlashOCC uses the same LSS frustum / bev-pool formulation in its camera-only occ branch.